

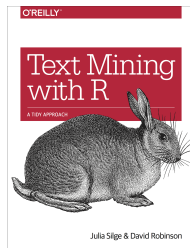
Introduction to Text Mining with R

RLadies Philadelphia



Today's "bookclub"

- Chapters 1 & 2 of *"Text Mining with R: A Tidy Approach"* by Julia Silge and David Robinson <https://www.tidytextmining.com/>
- Content overview
- Work together in small groups



The tidy text format

- Some keywords
 - **token** = meaningful unit of text
 - **tokenisation** = splitting text into tokens
 - **n-gram** = adjacent sequence of n items from sample of text
 - ▶ unigram, bigram, trigrams. . .
 - **regex** = “regular expression” = sequence of characters defining a search pattern
- What is “the tidy text format”?
 - one-token-per-row
 - dplyr, ggplot2

A minimal example

```
text <- c("Because I could not stop for Death -",  
         "He kindly stopped for me -",  
         "The Carriage held but just Ourselves -",  
         "and Immortality")  
text_df <- data_frame(line = 1:4, text = text)  
text_df
```

```
## # A tibble: 4 x 2  
##   line text  
##   <int> <chr>  
## 1     1 Because I could not stop for Death -  
## 2     2 He kindly stopped for me -  
## 3     3 The Carriage held but just Ourselves -  
## 4     4 and Immortality
```

Tokenisation

- `unnest_tokens()` from `tidytext` to tokenize
 - words (default)
 - characters
 - n-grams
 - sentences
 - lines
 - paragraphs
 - regex pattern separation

```
head(text_df %>%  
  unnest_tokens(word, text))
```

```
## # A tibble: 6 x 2  
##   line word  
##   <int> <chr>  
## 1     1 because  
## 2     1 i  
## 3     1 could  
## 4     1 not  
## 5     1 stop  
## 6     1 for
```

Processing austen_books()

- austen_books() from janeaustenr
- mutate to add
 - line numbers for each row
 - chapter (w/ cumulative sum of regex string finds)

```
original_books <- austen_books() %>%  
  group_by(book) %>%  
  mutate(linenumber = row_number(),  
         chapter = cumsum(str_detect(text,  
                                regex("^chapter [\\divxlc]",  
                                ignore_case = TRUE)))) %>%  
  ungroup()  
original_books[1,]
```

```
## # A tibble: 1 x 4  
##   text                book      linenumber chapter  
##   <chr>              <fct>          <int>    <int>  
## 1 SENSE AND SENSIBILITY Sense & Sensibility      1        0
```

Processing austen_books()

- tokenise words

```
tidy_books <- original_books %>%  
  unnest_tokens(word, text)  
head(tidy_books)
```

```
## # A tibble: 6 x 4  
##   book                linenumber chapter word  
##   <fct>                <int>    <int> <chr>  
## 1 Sense & Sensibility     1        0 sense  
## 2 Sense & Sensibility     1        0 and  
## 3 Sense & Sensibility     1        0 sensibility  
## 4 Sense & Sensibility     3        0 by  
## 5 Sense & Sensibility     3        0 jane  
## 6 Sense & Sensibility     3        0 austen
```

stop_words

- stop_words from tidytext
- purpose?
- remove stop_words with an antijoin

```
tidy_books <- tidy_books %>%  
  anti_join(stop_words)
```




The most frequent words

```
tidy_books %>%  
  count(word, sort = TRUE)
```

```
## # A tibble: 13,914 x 2  
##   word      n  
##   <chr> <int>  
## 1 miss    1855  
## 2 time    1337  
## 3 fanny    862  
## 4 dear     822  
## 5 lady     817  
## 6 sir      806  
## 7 day      797  
## 8 emma     787  
## 9 sister   727  
## 10 house   699  
## # ... with 13,904 more rows
```

The most frequent words



Project Gutenberg

- 1971
- <https://www.gutenberg.org/>
- 56,000 ebooks
- public domain / expired copyright

gutenbergr

- gutenbergr package
- https://ropensci.org/tutorials/gutenbergr_tutorial/
- Search `gutenberg_metadata`

```
head(gutenberg_metadata %>%  
  filter(author == "Wells, H. G. (Herbert George)"))  
head(gutenberg_works(author == "Wells, H. G. (Herbert George)"))
```

- Download by id

```
hgwells <- gutenberg_download(c(35, 36, 5230, 159))  
bronte <- gutenberg_download(c(1260, 768, 969, 9182, 767))
```

Word frequencies for Jane Austen, the Brontë sisters, and H.G. Wells



Sentiment analysis

- aka “opinion mining”
- computationally identifying & categorizing sentiment in text
- Some keywords
 - **lexicon** = inventory of words
 - **sentiment** = emotional content

The sentiments dataset (tidytext)

- contains several sentiment lexicons

```
unique(sentiments$lexicon)
```

```
## [1] "nrc"      "bing"     "AFINN"    "loughran"
```

- Limitations of this approach?
 - appropriateness
 - unigrams
 - ▶ *"It was not good"*

The lexicons

- **AFINN** (Finn Årup Nielsen)

- -5 to 5
- manually labelled

```
head(get_sentiments("afinn"))
```

```
## # A tibble: 6 x 2
##   word      score
##   <chr>    <int>
## 1 abandon      -2
## 2 abandoned    -2
## 3 abandons     -2
## 4 abducted     -2
## 5 abduction    -2
## 6 abductions   -2
```


The lexicons

- **bing** (Bing Liu et al)
 - positive/negative

```
head(get_sentiments("bing"))
```

```
## # A tibble: 6 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 2-faced   negative
## 2 2-faces   negative
## 3 a+       positive
## 4 abnormal  negative
## 5 abolish   negative
## 6 abominable negative
```

The lexicons

- **nrc** (Saif Mohammad and Peter Turney)
 - positive/negative
 - anger, fear, anticipation, trust, surprise, sadness, joy, disgust
 - crowdsourced manual annotations

```
head(get_sentiments("nrc"))
```

```
## # A tibble: 6 x 2
##   word      sentiment
##   <chr>     <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
## 4 abandon   sadness
## 5 abandoned anger
## 6 abandoned fear
```

Most frequent words by sentiment

- Most frequent “joy” words in Emma

```
nrc_joy <- get_sentiments("nrc") %>%  
  filter(sentiment == "joy")  
tidy_books %>%  
  filter(book == "Emma") %>%  
  inner_join(nrc_joy) %>%  
  count(word, sort = TRUE)
```

```
## # A tibble: 298 x 2  
##   word      n  
##   <chr>    <int>  
## 1 friend   166  
## 2 hope     143  
## 3 happy    125  
## 4 love     117  
## 5 deal      92  
## 6 found     92  
## 7 happiness  76  
## 8 pretty    68  
## 9 true      66  
## 10 comfort  65  
## # ... with 288 more rows
```


Sentiment change across text

- 80 line chunks

```
jane_austen_sentiment <- tidy_books %>%  
  inner_join(get_sentiments("bing")) %>%  
  count(book, index = linenumber %/% 80, sentiment) %>%  
  spread(sentiment, n, fill = 0) %>%  
  mutate(sentiment = positive - negative)  
jane_austen_sentiment
```

```
## # A tibble: 920 x 5  
##   book                                index negative positive sentiment  
##   <fct>                             <dbl>     <dbl>     <dbl>     <dbl>  
## 1 Sense & Sensibility                0.         16.         26.         10.  
## 2 Sense & Sensibility                1.         19.         44.         25.  
## 3 Sense & Sensibility                2.         12.         23.         11.  
## 4 Sense & Sensibility                3.         15.         22.          7.  
## 5 Sense & Sensibility                4.         16.         29.         13.  
## 6 Sense & Sensibility                5.         16.         39.         23.  
## 7 Sense & Sensibility                6.         24.         37.         13.  
## 8 Sense & Sensibility                7.         22.         39.         17.  
## 9 Sense & Sensibility                8.         30.         35.          5.  
## 10 Sense & Sensibility               9.         14.         18.          4.  
## # ... with 910 more rows
```

Sentiment change across Austen novels



Most frequent positive versus negative words

```
bing_word_counts <- tidy_books %>%  
  inner_join(get_sentiments("bing")) %>%  
  count(word, sentiment, sort = TRUE)  
  
head(bing_word_counts)
```

```
## # A tibble: 6 x 3  
##   word      sentiment      n  
##   <chr>    <chr>    <int>  
## 1 miss      negative  1855  
## 2 happy     positive   534  
## 3 love      positive   495  
## 4 pleasure positive   462  
## 5 poor      negative   424  
## 6 happiness positive   369
```

Custom stop words

- “miss” is currently misanalysed...

```
custom_stop_words <- bind_rows(data_frame(word = c("miss"),  
                                           lexicon = c("custom")),  
                               stop_words)  
  
head(custom_stop_words)
```

```
## # A tibble: 6 x 2  
##   word  lexicon  
##   <chr> <chr>  
## 1 miss  custom  
## 2 a     SMART  
## 3 a's   SMART  
## 4 able  SMART  
## 5 about SMART  
## 6 above SMART
```

Pipe data into a wordcloud plot...

- wordcloud package

```
tidy_books %>%
  anti_join(custom_stop_words) %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  acast(word ~ sentiment, value.var = "n", fill = 0) %>%
  comparison.cloud(colors = c("red", "blue"),
                  max.words = 100)
```



In small groups

- Work through Section 1.5 (Chapter 1), especially the code chunk below:

```
frequency <- bind_rows(mutate(tidy_bronte, author = "Bronte Sisters"),
                        mutate(tidy_hgwells, author = "H.G. Wells"),
                        mutate(tidy_books, author = "Jane Austen")) %>%
  mutate(word = str_extract(word, "[a-z']+")) %>%
  count(author, word) %>%
  group_by(author) %>%
  mutate(proportion = n / sum(n)) %>%
  select(-n) %>%
  spread(author, proportion) %>%
  gather(author, proportion, `Bronte Sisters`:`H.G. Wells`)
```

- Pick a novel from the Gutenberg collection, and discover...
 - the most frequent words
 - the most frequent “trust” words using the nrc lexicon
 - how the sentiment changes across the novel using the bing lexicon
- If you have time, create a data visualisation for one of these text mining exercises

References

- Silge, Julia and David Robinson (2017), *Text Mining with R: A Tidy Approach*, O'Reilly / <https://www.tidytextmining.com>
 - and references therein:
<https://www.tidytextmining.com/references.html>

Dracula: the most frequent words

```
gutenberg_works(title == "Dracula")$gutenberg_id
```

```
## [1] 345
```

```
dracula <- gutenberg_download(345)
dracula_words <- dracula %>%
  mutate(linenummer = row_number(),
         chapter = cumsum(str_detect(text,
                                     regex("^chapter [\\divxlc]",
                                     ignore_case = TRUE)))) %>%
  unnest_tokens(word, text)
head(dracula_words %>%
  anti_join(stop_words) %>%
  count(word, sort = TRUE))
```

```
## # A tibble: 6 x 2
##   word      n
##   <chr>   <int>
## 1 time     390
## 2 van      323
## 3 night    310
## 4 helsing  301
## 5 dear     224
## 6 lucy     223
```


Dracula: the most frequent “trust” words using the nrc lexicon

```
nrc_trust <- get_sentiments("nrc") %>%  
  filter(sentiment == "trust")  
dracula_words %>%  
  inner_join(nrc_trust) %>%  
  count(word, sort = TRUE)
```

```
## # A tibble: 471 x 2  
##   word      n  
##   <chr>    <int>  
## 1 good      258  
## 2 friend    184  
## 3 professor 155  
## 4 count     153  
## 5 found     153  
## 6 god       150  
## 7 diary     110  
## 8 white     107  
## 9 lord       79  
## 10 hope      66  
## # ... with 461 more rows
```

Dracula: how the sentiment changes across the novel using the bing lexicon

```
dracula_sentiment <- dracula_words %>%  
  inner_join(get_sentiments("bing")) %>%  
  count(index = linenumbers%%80, sentiment) %>%  
  spread(sentiment, n,  
    fill = 0) %>% mutate(sentiment = positive - negative)  
ggplot(dracula_sentiment, aes(index, sentiment, fill=sentiment)) +  
  scale_fill_gradient(low = "#230105", high = "#ff0019") +  
  geom_col(show.legend = FALSE)
```



Dracula: a wordcloud

```
red_palette <- brewer.pal(8,"RdGy")
dracula_words %>%
  anti_join(stop_words) %>%
  count(word) %>%
  with(wordcloud(word, n, min.freq=6, max.words = 250,
    random.order=FALSE, rot.per=.15, colors=red_palette))
```



Work through each line in the code chunk below from Section 1.5:

- h/t Alice!
 - `bind_rows` is like `rbind()` but in `dplyr`, add an author variable to each tibble
 - `mutate` is changing the word variable by looking for the regex starting with lower case letters
 - counting words by author
 - `dplyr group_by` gives each observation a group and performs next functions on those groups
 - `mutate` creates proportion variable (count of each word divided by the total counts)
 - `select(-n)` removes the `n` (count) variable from the tibble
 - `spread()` is creating a wide format table where each author has a variable
 - `gather()` is pulling back together so Austen has a variable, but Bronte and Wells are in long format

Packages

```
library(dplyr)
library(tidytext)
library(janeaustenr)
library(dplyr)
library(stringr)
library(ggplot2)
library(gutenbergr)
library(tidyr)
library(scales)
library(wordcloud)
library(reshape2)
```


Potential problems

- Problem loading wordcloud with R version 3.3.3
- <https://stackoverflow.com/questions/39885408/dependency-slam-is-not-available-when-installing-tm-package>
- Probably better to update R?

```
# install.packages('devtools')  
# library(devtools)  
# slam_url <- "https://cran.r-project.org/src/contrib/...  
# ...Archive/slam/slam_0.1-37.tar.gz"  
# install_url(slam_url)
```

Potential problems

- “Error in summarise_impl(.data, dots) : invalid argument type”
- Update dplyr

```
head(data_frame(text = prideprejudice) %>%  
  unnest_tokens(sentence, text, token = "sentences"))
```

```
## # A tibble: 6 x 1  
##   sentence  
##   <chr>  
## 1 pride and prejudice by jane austen chapter 1 it is a truth univer~  
## 2 however little known the feelings or views of such a man may be on his ~  
## 3 "\"my dear mr."  
## 4 "bennet,\" said his lady to him one day, \"have you heard that netherfi~  
## 5 mr.  
## 6 bennet replied that he had not.
```